

Kokkos

A Parallel, Portable Programming Model for CPUs and GPUs

Rohit Kakodkar

Research Software Engineer



What is Kokkos?

- Kokkos is a C++^{*} performance portability library
 - Write a single source implementation
 - Descriptive programming model
 - Compile for CPUs or GPUs
 - Kokkos is a shared memory programming model (works in conjunction with MPI)
- Major buy-in by DOE and national labs
 - LAAMPS, Trilinos, PETSc
 - Over 100 projects using kokkos
 - Contributing to the C++ standard
 - Active kokkos developers community via slack (invite only)





Python and Fortran bindings are available

The goal here is to motivate a use case for Kokkos. Given time limit this is not an in-depth tutorial of Kokkos (probably sometime in the future).

- Outline of the talk:
 - Need for performance portability
 - Understanding performance
 - Caching (CPUs) vs Coalescing (GPUs)
 - Kokkos views
 - Kokkos parallel construct
 - Hands-on example
 - Parallel policies
 - Packaging Kokkos projects
 - Why use kokkos?

Need for performance portability

Modern high-performance computing applications need to run efficiently across multiple architectures:

		GPU		CPU
Architecture	 			
Programming model	CUDA	HIP	DPC++ (SYCL)	OpenMP pthreads

Need for performance portability

Modern high-performance computing applications need to run efficiently across multiple architectures:

Pre-Exascale



LANL Trinity
Intel Haswell/ Intel KNL
OpenMP 3



ORNL Summit
NVIDIA Volta100
OpenMP/CUDA

Exascale



ORNL Frontier
AMD GPUs
OpenMP/HIP



ANL Aurora
Intel GPUs
DPC++

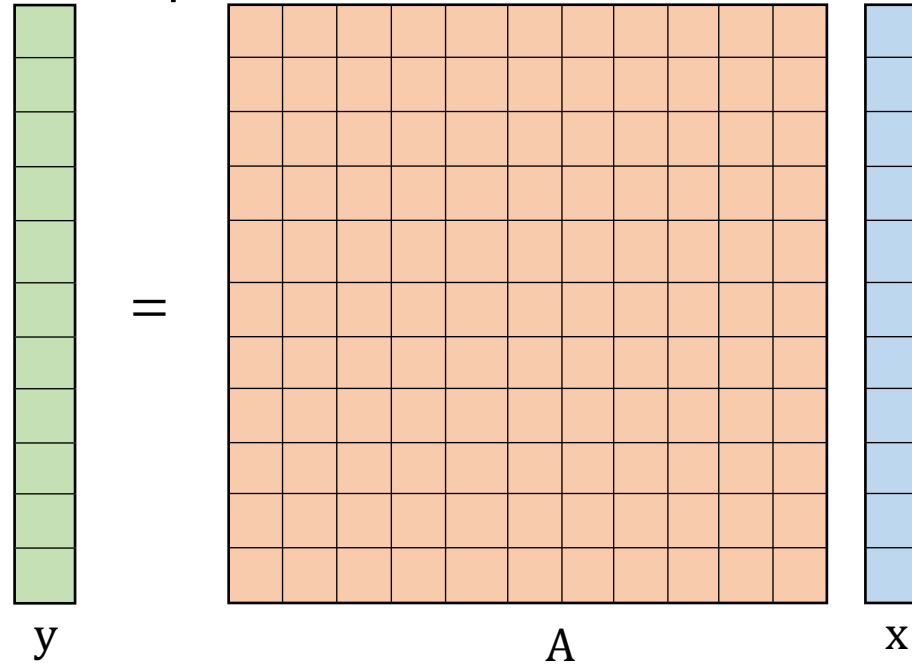
Moving into the exascale era will require adapting your applications to utilize modern architectures

Need for performance portability

- Problem: Porting applications for various architectures is time consuming
 - Typical HPC application : 300k – 600k lines of code
 - Smaller scale applications : SpecFEM2D + SpecFEM3D ~ 100k lines of code (conservative estimate)
 - Porting requires ~10% rewrite of the application
 - Typical software engineer writes about 20k LOC/year
- Potential portability options: OpenMP 5 or OpenACC

What about performance??

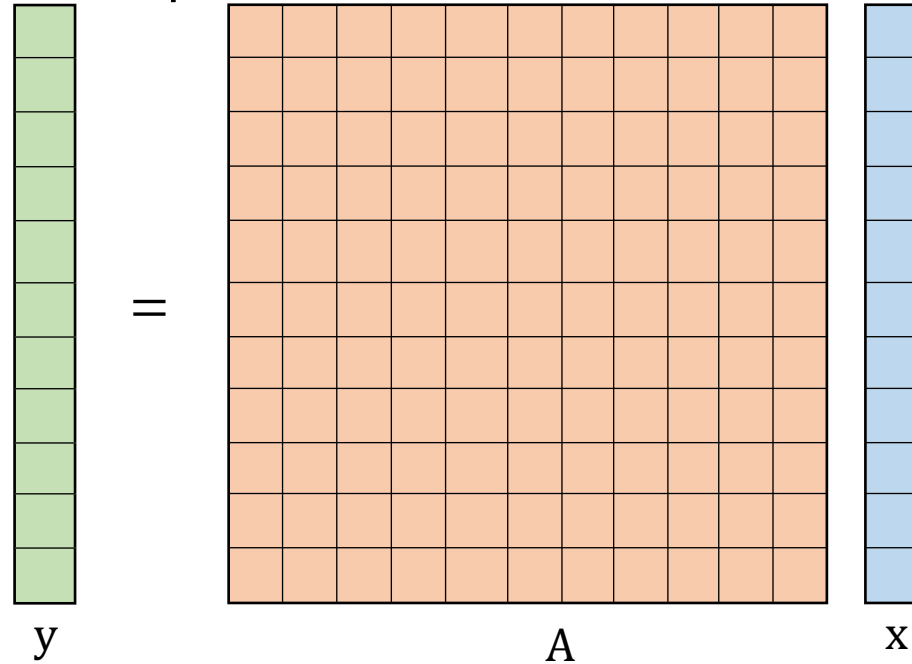
- Example: Matrix – vector multiplication



Serial Implementation

```
for (int j = 0; j < N; j++){  
    for (int i = 0; i < M; i++){  
        y(j) += A(j,i)*x(i);  
    }  
}
```

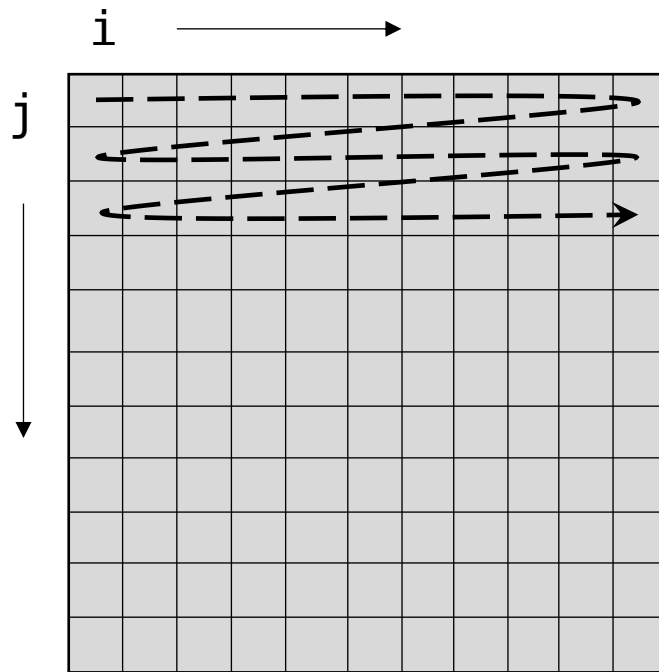
- Example: Matrix – vector multiplication



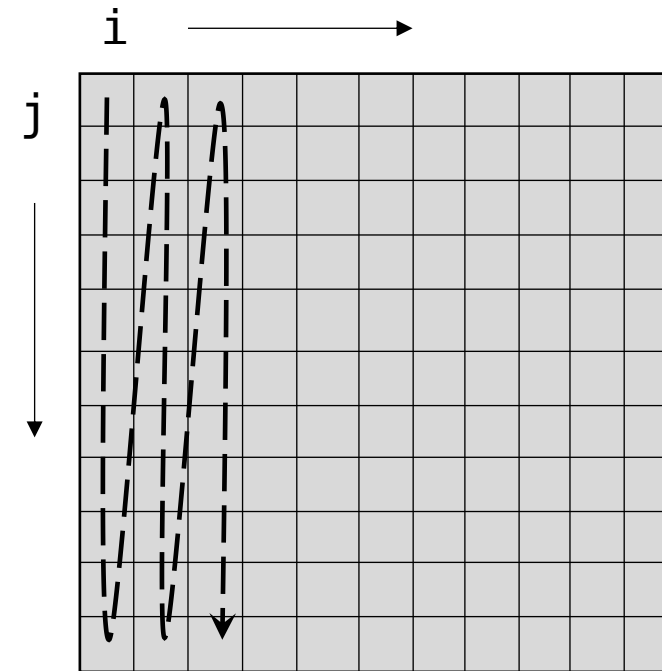
Parallel Implementation

```
parallel_for (int j = 0; j < N; j++){  
    // distribute different rows to different threads  
    for (int i = 0; i < M; i++){  
        y(j) += A(j,i)*x(i);  
    }  
}
```


- Performance is dependent on matrix layout

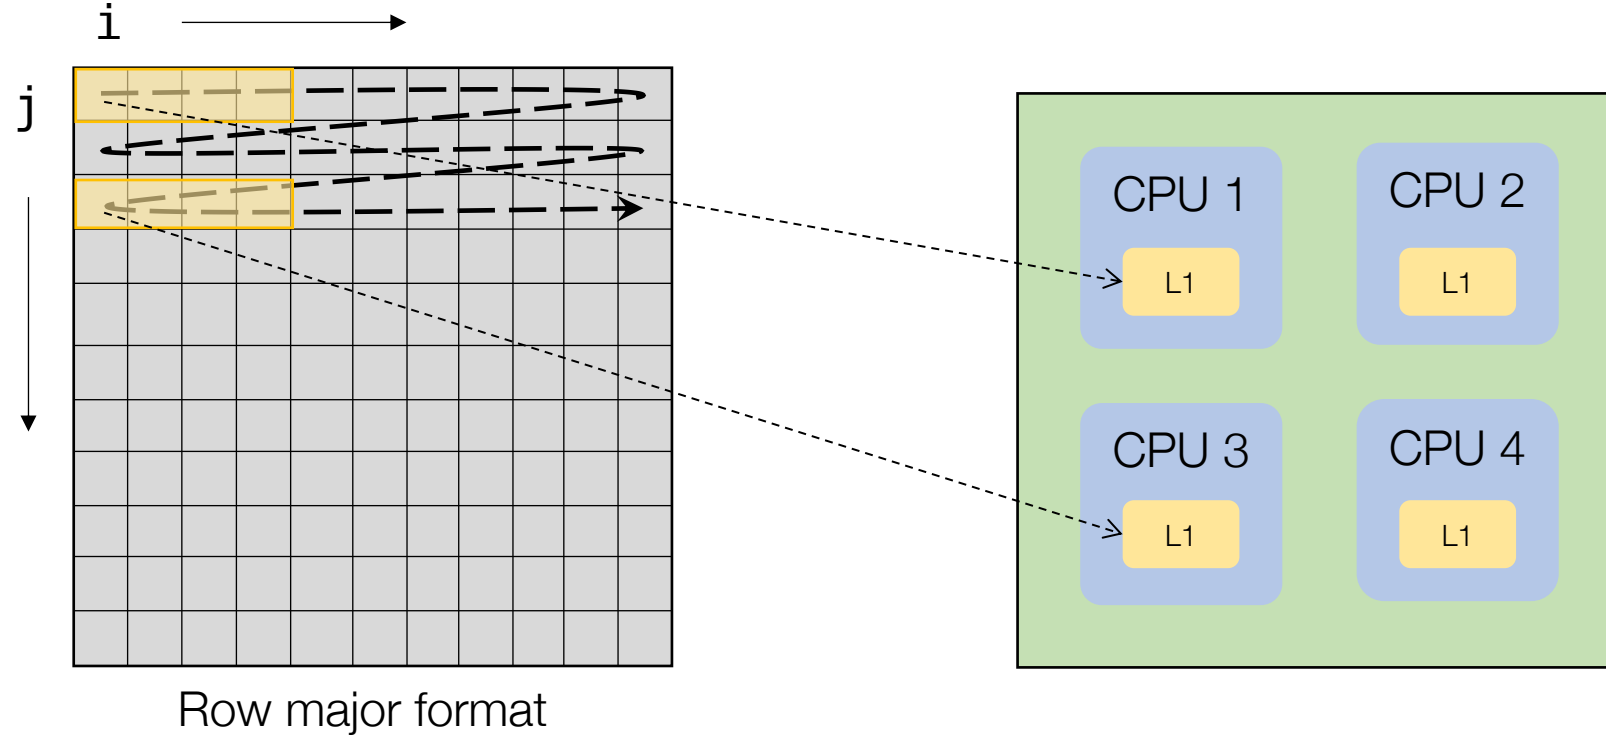


Row major format



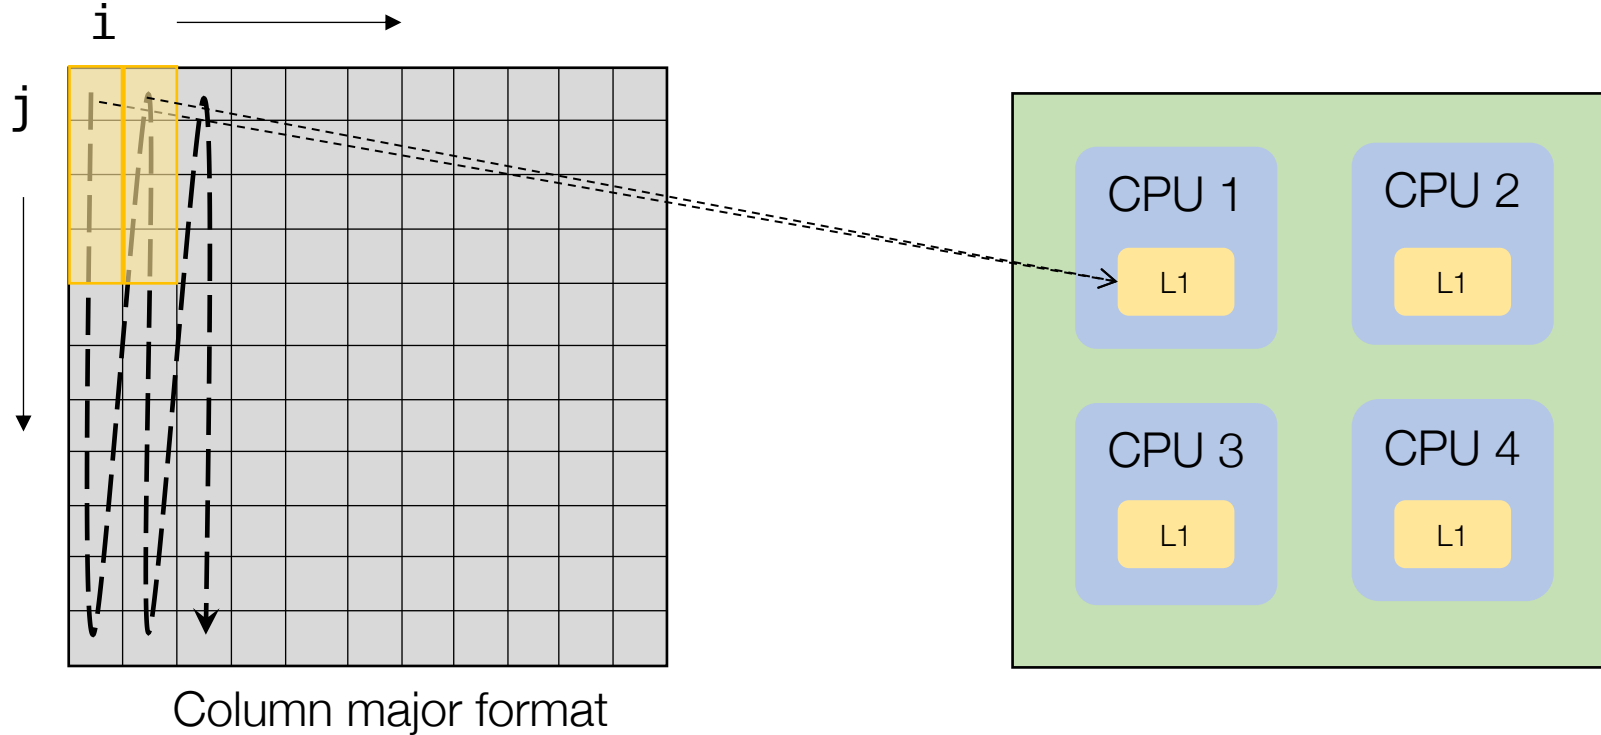
Column major format

- Good performance



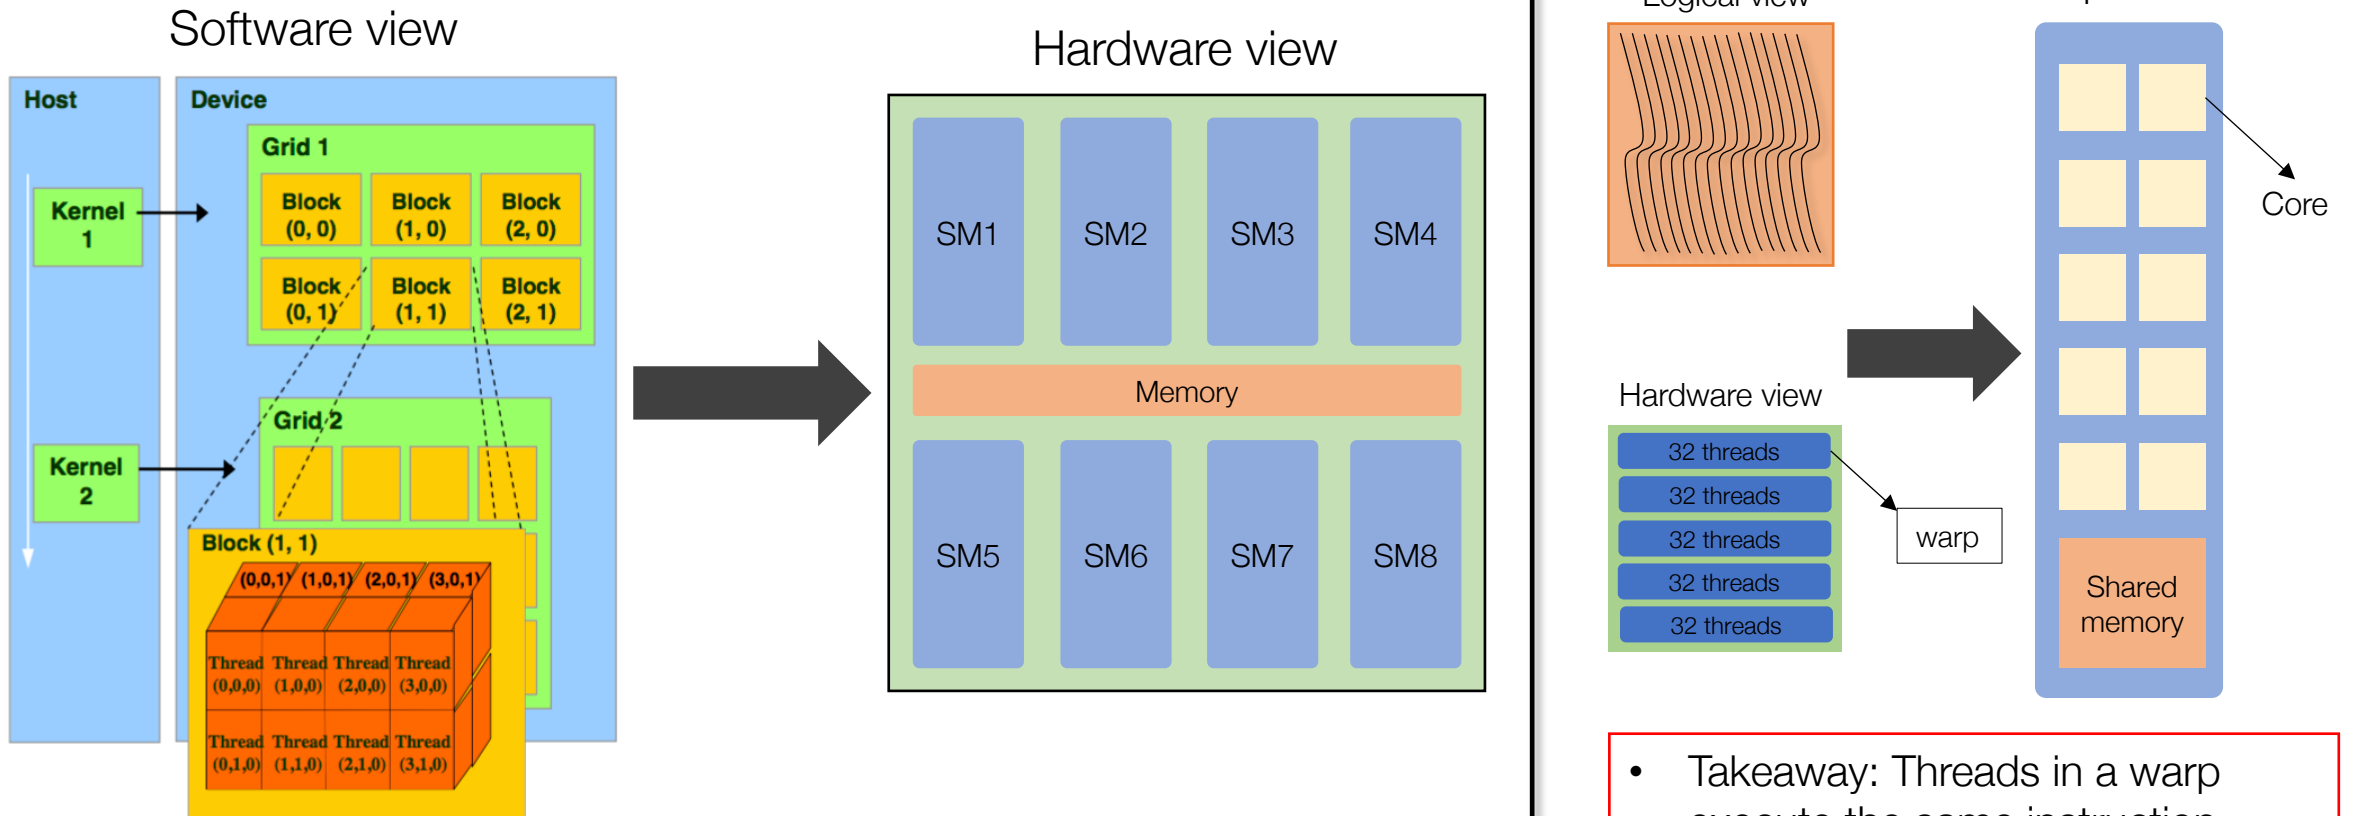
```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```

- Bad performance



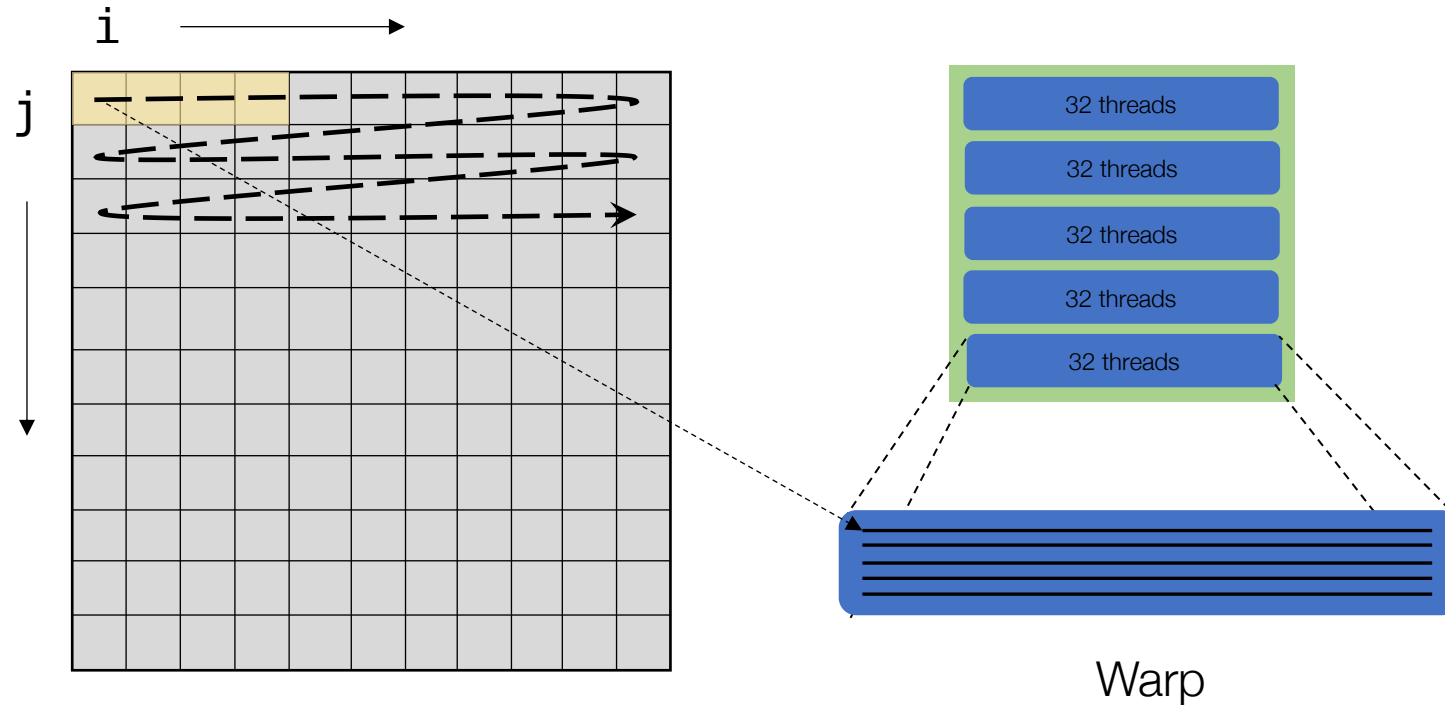
```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```

- GPU execution model



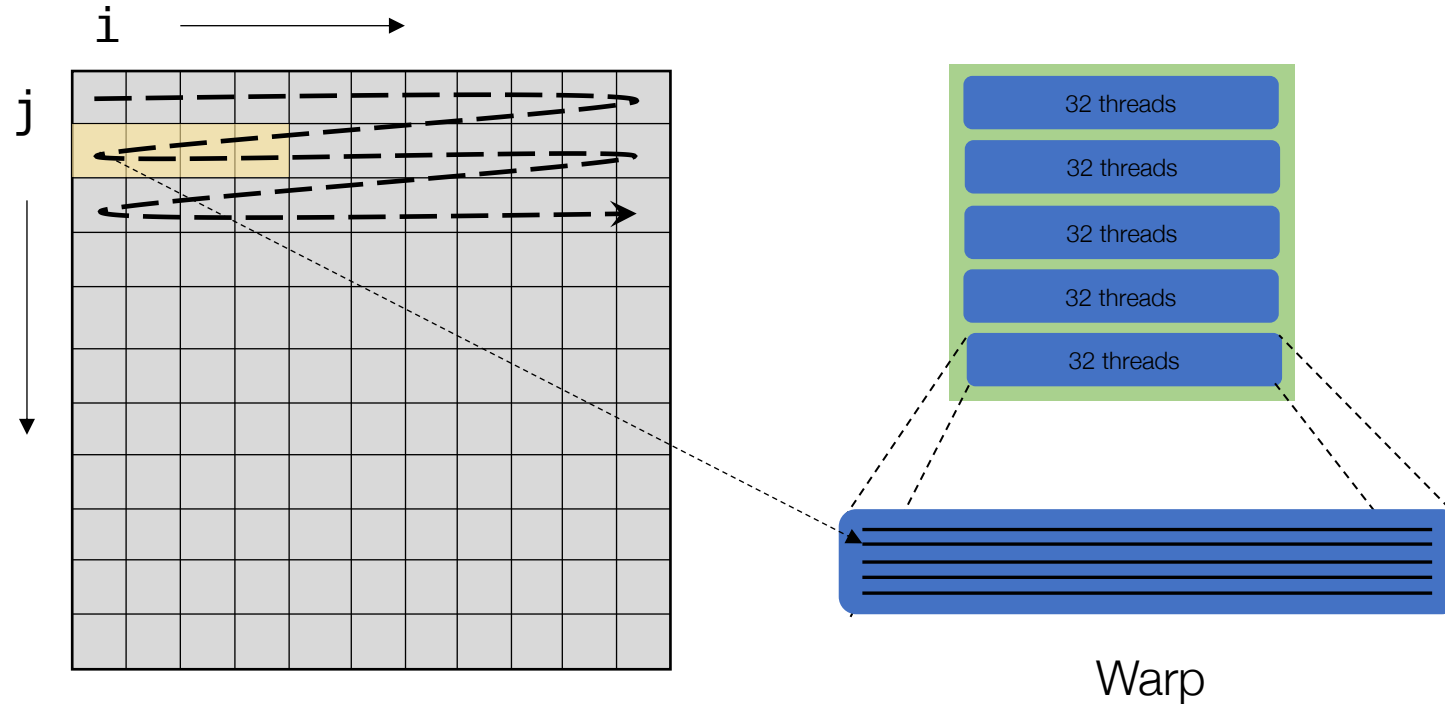
- Takeaway: Threads in a warp execute the same instruction

- Bad performance



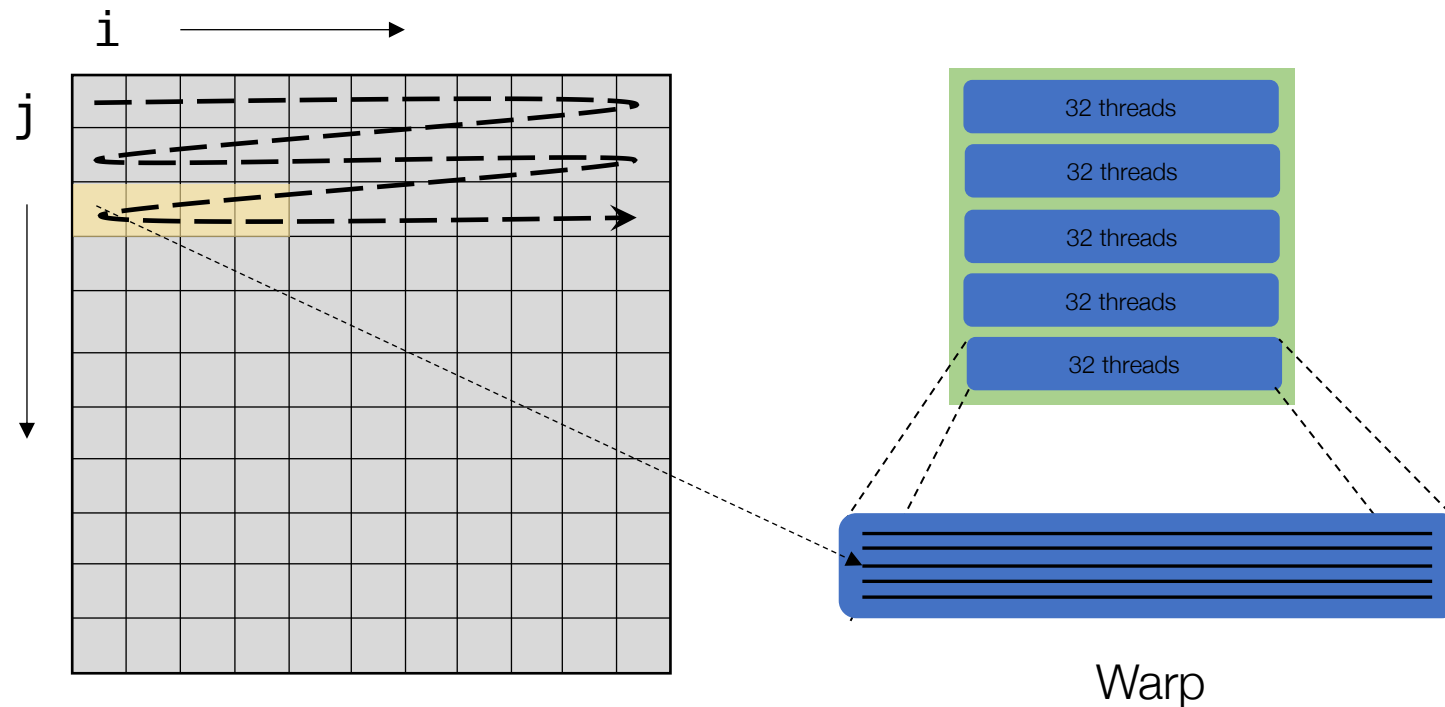
```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```

- Bad performance



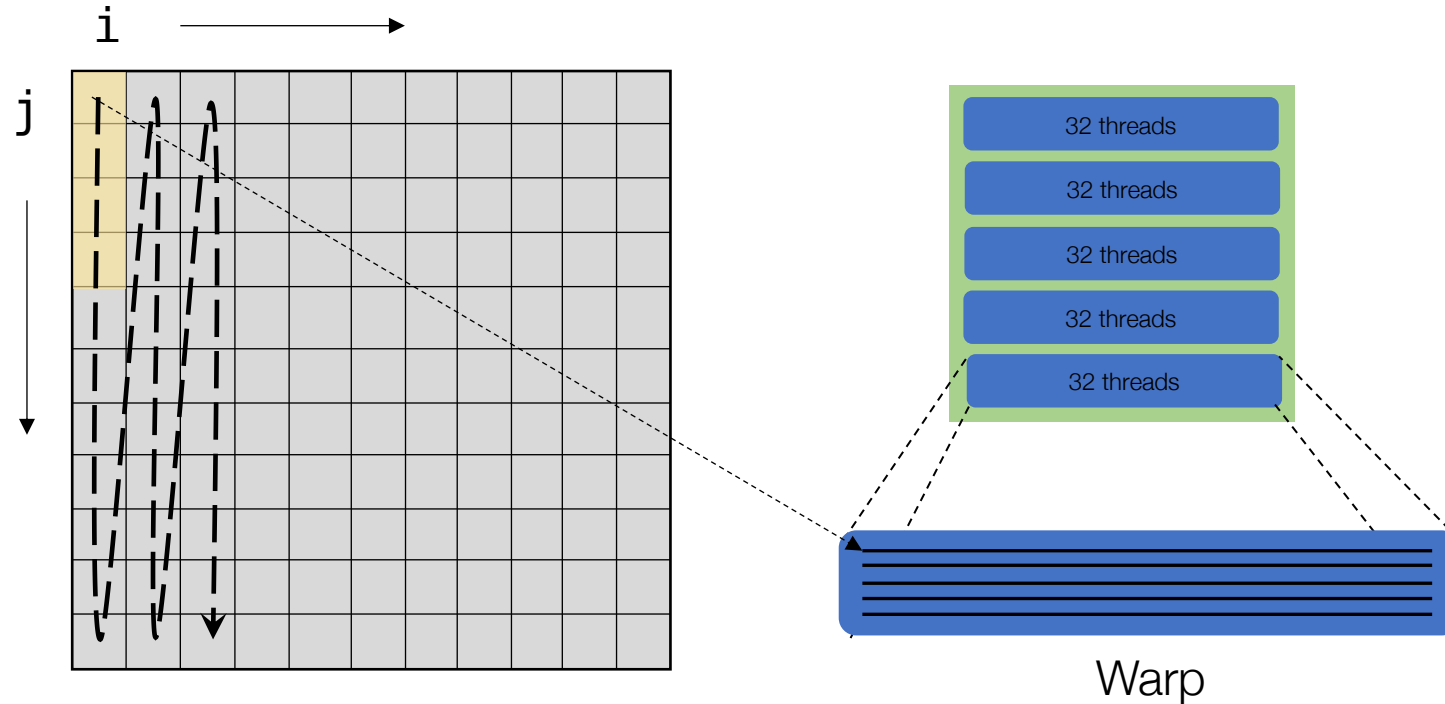
```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```

- Bad performance



```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```

- Good performance

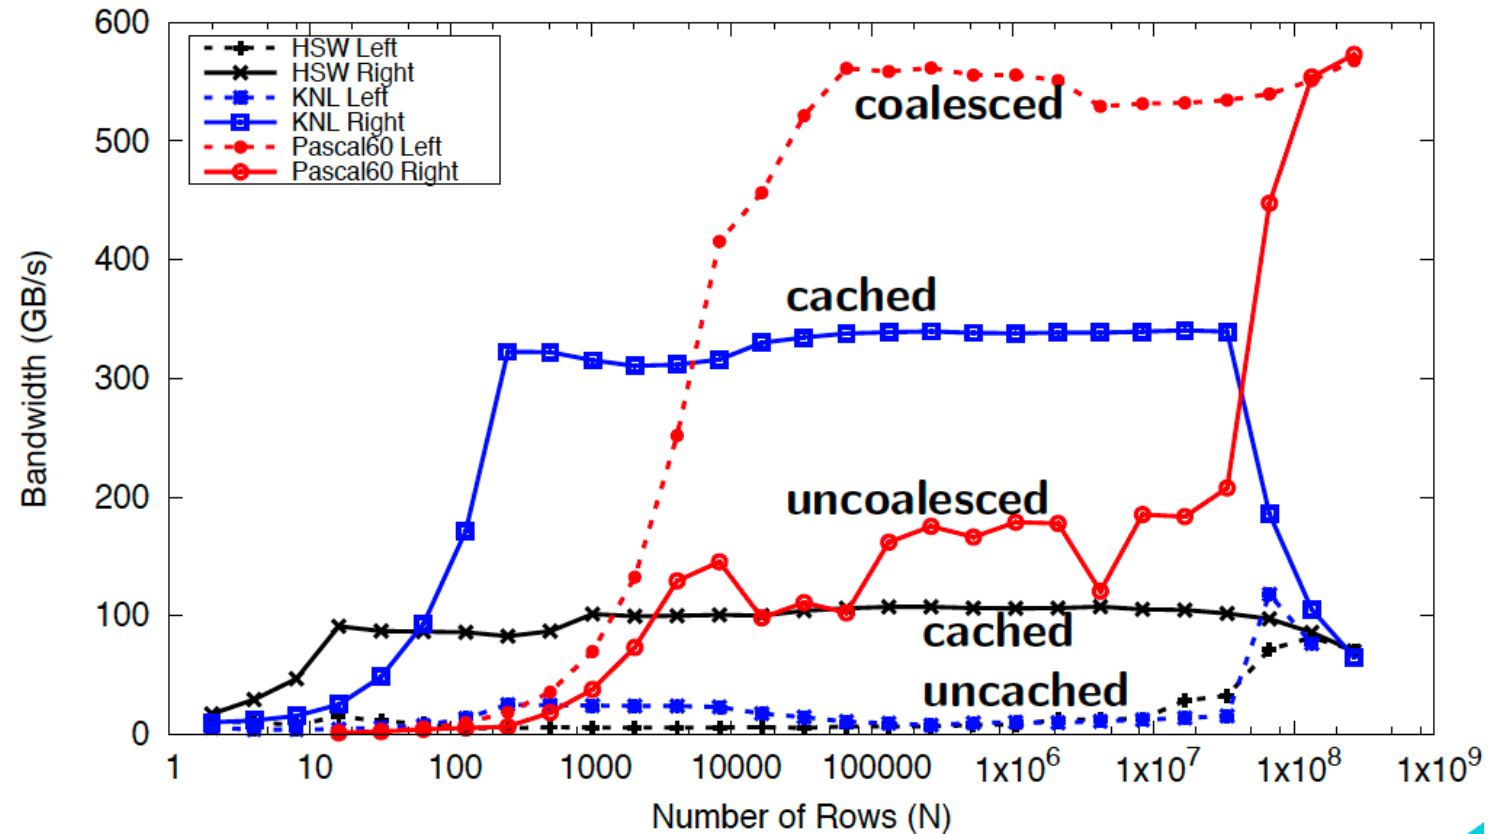


```
parallel_for j: N
// every thread gets a specific j to execute
for i: M
    A(j,i)
```


- CPUs require row major format
 - $A(i_1, i_2, \dots, i_n) \rightarrow i_n$ needs to be continuous in memory (Layout Right)
 - Results in data caching
- GPUs require a column major format
 - $A(i_1, i_2, \dots, i_n) \rightarrow i_1$ needs to be continuous in memory (Layout Left)
 - Results in data coalescing

Performance benchmarks

- vector - matrix – vector multiplication
- Total number of elements in matrix is a constant ($M \times N = \text{constant}$)
- Bandwidth \sim amount of data processed per second



- CPUs require row major format
 - $A(i_1, i_2, \dots, i_n) \rightarrow i_n$ needs to be continuous in memory (Layout Right)
 - Results in data caching
- GPUs require a column major format
 - $A(i_1, i_2, \dots, i_n) \rightarrow i_1$ needs to be continuous in memory (Layout Left)
 - Results in data coalescing

Data layout needs to be determined at compile time based on target architecture

- Kokkos views are multidimensional arrays
- Can reside on CPUs or GPUs
 - Use “mirror” of your arrays to sync between CPUs and GPUs
 - `deep_copy` can be initialized between original and mirror views
- View layout can be specified at compile time
 - Default behavior is to optimized layout based on execution space
 - GPUs = Layout Left
 - CPUs = Layout Right
- Views are reference counted
 - Easy memory management
- Copy construction leads to both views pointing to the same data
 - deep copy requires use of “`deep_copy`” routine rather than a copy construction

```
Kokkos::View<float **, Kokkos::LayoutLeft, Kokkos::CudaSpace>  
A("A_matrix", N, M);
```

Policy

Pattern `for` (`int element = 0; element < nelements; element++`){
 `total = 0;`
 `for (int qp = 0; qp < numQP; qp++){`
 `// reduction over every qp in element`
 `total += foo(variable, element, qp);`
 `}`
 `elementValue[element] = total;`
}

Body

OpenMP – An analogy

Policy

Pattern

```
#pragma omp parallel for schedule(<static, dynamic>)  
for (int element = 0; element < nelements; element++){  
    total = 0;  
    for (int qp = 0; qp < numQP; qp++){  
        // reduction over every qp in element  
        total += foo(variable, element, qp);  
    }  
    elementValue[element] = total;  
}
```

Body

Pattern

Policy

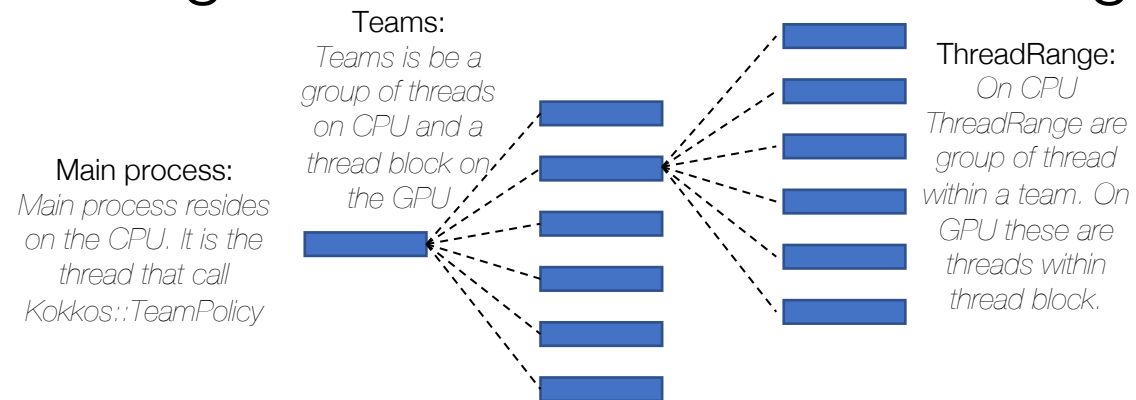
```
Kokkos::parallel_for("LoopName", Kokkos::RangePolicy(nelements),  
    [=](const int element){  
        total = 0;  
        for (int qp = 0; qp < numQP; qp++){  
            // reduction over every qp in element  
            total += foo(variable, element, qp);  
        }  
        elementValue[element] = total;  
    });
```

Body

Example

Parallel policies

- `Kokkos::RangePolicy` – Each iterate is an integer in a contiguous range
- `Kokkos::MDRangePolicy` – Each iterate for each rank is an integer in a contiguous range
- `Kokkos::TeamPolicy` – Assigns to each iterate in a contiguous range a team of threads



Packaging Kokkos projects

- CMAKE - Recommended
 - Provide Kokkos as a dependency package

```
include(FetchContent)
FetchContent_Declare(
  kokkos
  URL
  https://github.com/kokkos/kokkos/archive/refs/tags/3.6.01.zip
)
FetchContent_MakeAvailable(kokkos)

add_executable(question question.cpp)
target_link_libraries(question Kokkos::kokkos)
```

- CMAKE – Recommended
 - In house installations

```
set(KOKKOS_PATH <PATH TO KOKKOS ROOT>)  
  
add_subdirectory(KOKKOS_PATH)  
add_executable(question question.cpp)  
target_link_libraries(question Kokkos::kokkos)
```

- Makefiles – maybe suitable for individual private projects
 - Check this example out:
https://github.com/kokkos/kokkos/blob/master/benchmarks/bytes_and_flops/Makefile

Documentation Overview

Why use Kokkos?

	CUDA/HIP/DPC++	OpenMP 5/ OpenACC	Kokkos
Portability across architectures	No. Need to write separate kernels for every architecture	Yes. Single source code with pragma-based approach	Yes. Single source code implemented using Kokkos functions
Performance	Optimized performance	Tough to optimize for performance	Very good performance
Cost of portability	Very high	Medium	High
Cost of maintenance	Very high. Newer architectures might require tuning of kernels	Low. Assuming compilers do a good job of implementing the standard	Low. Assuming Kokkos backend is always optimized
Compiler dependence	N/A	These are standards, vendors have a flexibility on implementation	N/A
Fortran support	No. Could use bindings	Yes	No. Could use bindings

Why use Kokkos?

- Kokkos is a ***performance portability*** library
- Performance is dependent on data layout
 - Kokkos Views are easy way to manage data layout
- Compile time definition of data using C++ template meta-programming
- Low maintenance overhead for future architectures

Useful resources:

- [Documentation](#)
- [Tutorials](#)
- [Slack](#) (invite only) : ctrott [a] sandia.gov
- Local help - cses@princeton.edu

Thank you